



УНИВЕРСИТЕТ ИТМО


Parallel algorithms for the analysis and synthesis of data

Sokhin Timur, PhD student

Intro

The history of parallel computing can be traced back to the 1960s

- We can mention D825 Modular Data Processing System (1962) as a SMP or ILLIAC IV (1964) as MPP systems.
- 1966 - Flynn creates a taxonomy of architecture (SISD, SIMD, MISD, MIMD)




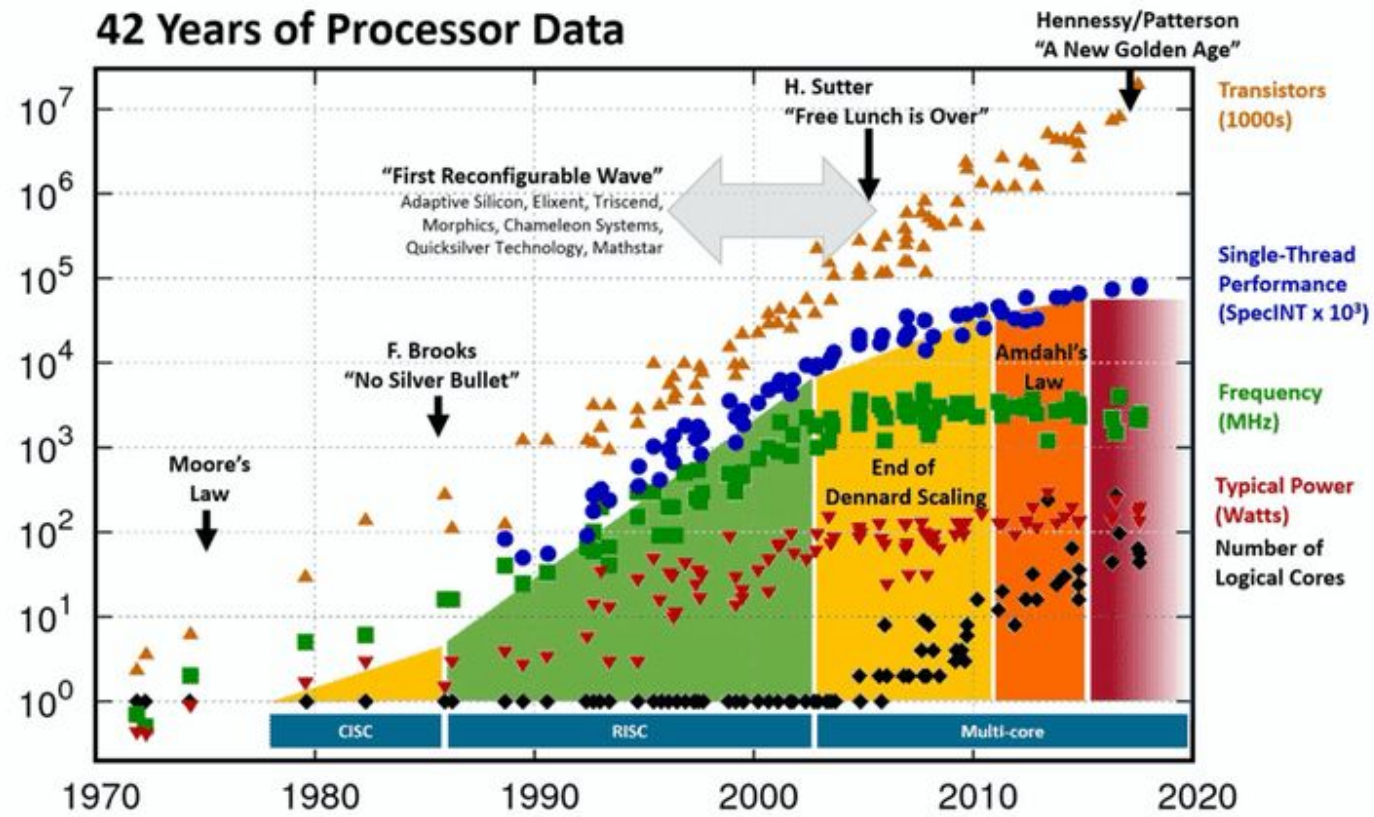
https://en.wikipedia.org/wiki/ILLIAC_IV#/media/File:ILLIAC_4_parallel_computer.jpg

- Cray-1 - example of vector processor based computer
- First mentioning of remote procedure call (RPC) in 1978 by Per Brinch Hansen
- 1982 - Cray X-MP with shared memory architecture




- 1993-94 MPI protocol
- 1997 OpenMP protocol
- 2007 CUDA platform






Hennessy and Patterson, Turing Lecture 2018, overlaid over "42 Years of Processors Data"
<https://www.karlsruhp.net/2018/02/42-years-of-microprocessor-trend-data/>; "First Wave" added by Les Wilson, Frank Schirrmeister
Original data up to the year 2010 collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond, and C. Batten
New plot and data collected for 2010-2017 by K. Rupp



Performance

Performance doesn't just depend on computing power and degree of parallelism - don't forget about how memory, networking mechanisms, etc. function.



Locality of reference

Temporal locality: If at one point a particular memory location is referenced, then it is likely that the same location will be referenced again in the near future.

Spatial locality: If a particular storage location is referenced at a particular time, then it is likely that nearby memory locations will be referenced in the near future.



```
for i in 0..n
  for j in 0..m
    for k in 0..p
      C[i][j] = C[i][j] + A[i][k] * B[k][j];
```

```
for i in 0..n
  for k in 0..m
    for j in 0..p
      C[i][j] = C[i][j] + A[i][k] * B[k][j];
```

Locality of reference



УНИВЕРСИТЕТ ИТМО

$$\begin{pmatrix} a_{00} & a_{01} & a_{02} \\ a_{10} & a_{11} & a_{12} \\ a_{20} & a_{21} & a_{22} \end{pmatrix} * \begin{pmatrix} b_{00} & b_{01} & b_{02} \\ b_{10} & b_{11} & b_{12} \\ b_{20} & b_{21} & b_{22} \end{pmatrix} = \begin{pmatrix} c_{00} & c_{01} & c_{02} \\ c_{10} & c_{11} & c_{12} \\ c_{20} & c_{21} & c_{22} \end{pmatrix}$$

```
for i in 0..n
```

```
  for j in 0..m
```

```
    for k in 0..p
```

```
      c[i][j] = c[i][j] + a[i][k] * b[k][j];
```

```
i=0, j=0, k=0
```

```
c[0][0] = c[0][0] + a[0][0] * b[0][0];
```

```
i=0, j=0, k=1
```

```
c[0][0] = c[0][0] + a[0][1] * b[1][0];
```

Memory

$[a_{00}, a_{01}, a_{02}, a_{10}, a_{11}, a_{12}, a_{20}, a_{21}, a_{22}]$

$[b_{00}, b_{01}, b_{02}, b_{10}, b_{11}, b_{12}, b_{20}, b_{21}, b_{22}]$

$[c_{00}, c_{01}, c_{02}, c_{10}, c_{11}, c_{12}, c_{20}, c_{21}, c_{22}]$

Registers

$[a_{00}, a_{01}, a_{02}, b_{00}, b_{01}, b_{02}, c_{00}, c_{01}, c_{02}]$

Locality of reference



УНИВЕРСИТЕТ ИТМО

$$\begin{pmatrix} a_{00} & a_{01} & a_{02} \\ a_{10} & a_{11} & a_{12} \\ a_{20} & a_{21} & a_{22} \end{pmatrix} * \begin{pmatrix} b_{00} & b_{01} & b_{02} \\ b_{10} & b_{11} & b_{12} \\ b_{20} & b_{21} & b_{22} \end{pmatrix} = \begin{pmatrix} c_{00} & c_{01} & c_{02} \\ c_{10} & c_{11} & c_{12} \\ c_{20} & c_{21} & c_{22} \end{pmatrix}$$

```
for i in 0..n
```

```
  for k in 0..m
```

```
    for j in 0..p
```

```
      c[i][j] = c[i][j] + a[i][k] * b[k][j];
```

```
i=0, j=0, k=0
```

```
c[0][0] = c[0][0] + a[0][0] * b[0][0];
```

```
i=0, j=1, k=0
```

```
c[0][1] = c[0][1] + a[0][0] * b[0][1];
```

Memory

$[a_{00}, a_{01}, a_{02}, a_{10}, a_{11}, a_{12}, a_{20}, a_{21}, a_{22}]$

$[b_{00}, b_{01}, b_{02}, b_{10}, b_{11}, b_{12}, b_{20}, b_{21}, b_{22}]$

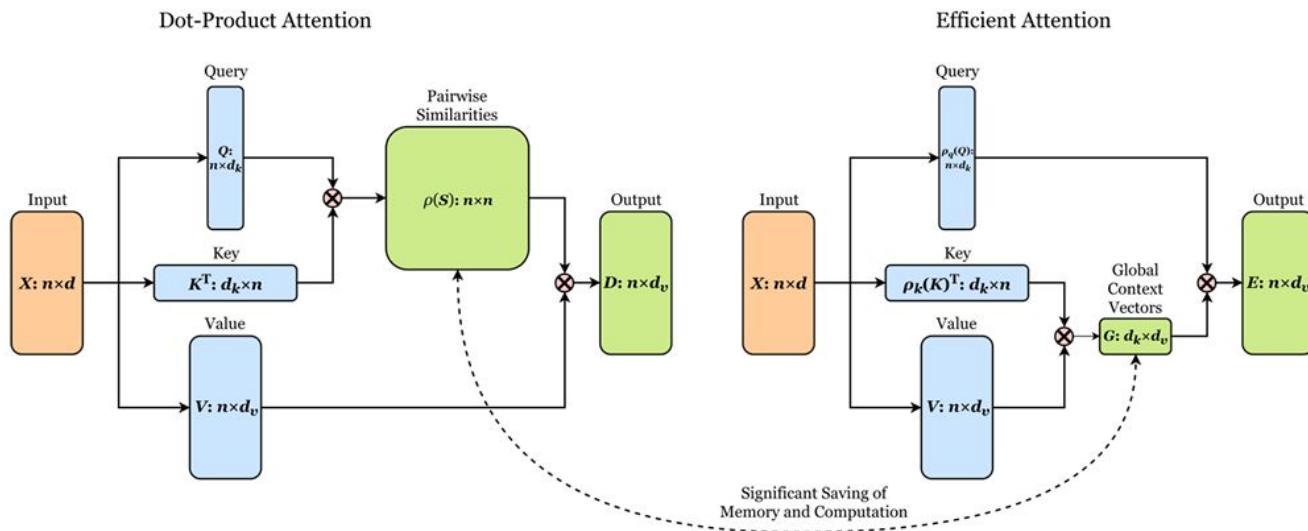
$[c_{00}, c_{01}, c_{02}, c_{10}, c_{11}, c_{12}, c_{20}, c_{21}, c_{22}]$

Registers

$[a_{00}, a_{01}, a_{02}, b_{00}, b_{01}, b_{02}, c_{00}, c_{01}, c_{02}]$

Computational complexity

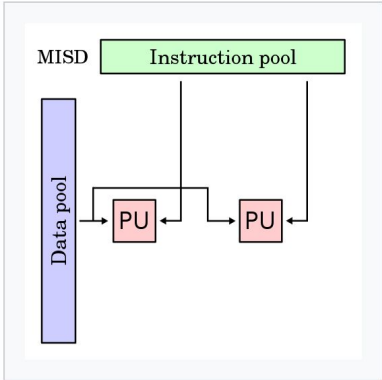

Transformer networks are well-parallelised, they have an ability to use buffers for inference





Flynn's taxonomy

single instruction multiple instruction

single data

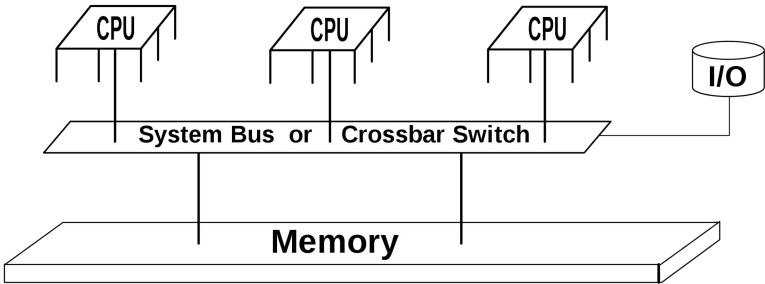


multiple data




Systematization by access to memory

Shared memory

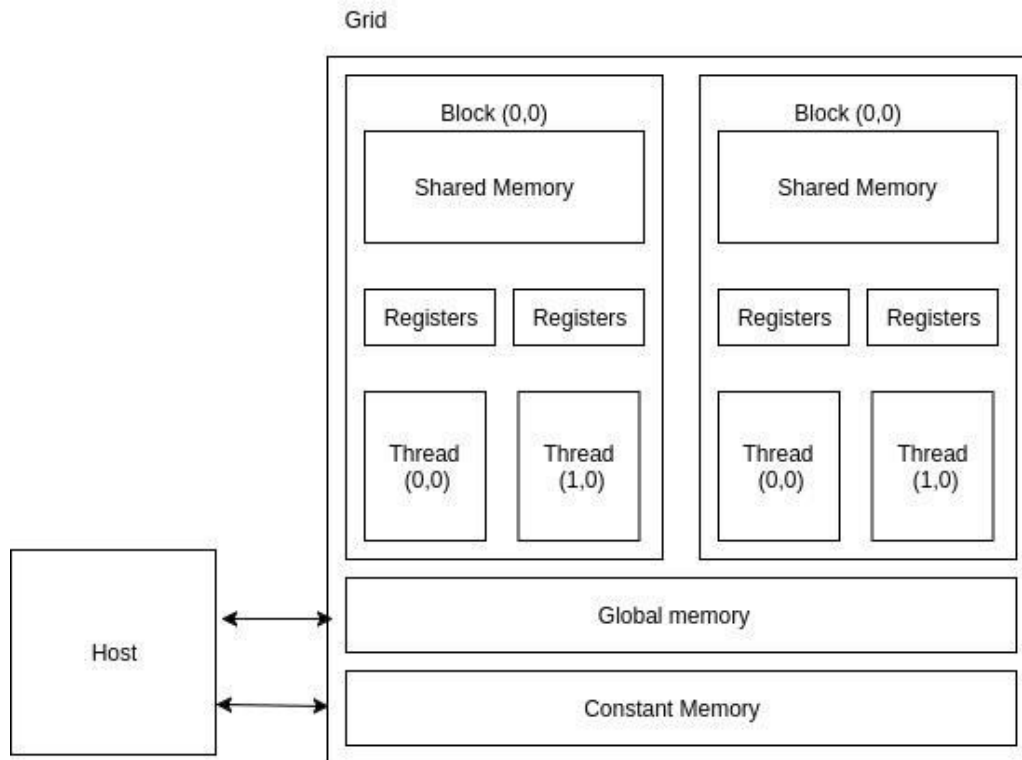


Distributed memory




Systematization by access to memory

CUDA memory model



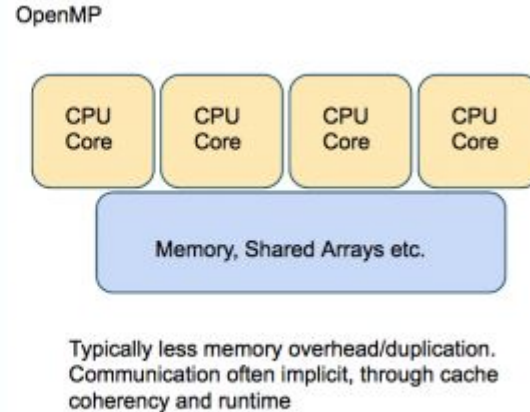
Parallel applications realization

Distributed memory




MPI (Message Passing Interface) is a standardized and portable message-passing standard designed to function on parallel computing architectures.

Shared memory



OpenMP (Open Multi-Processing) is an application programming interface (API) that supports multi-platform shared-memory multiprocessing programming.

Process



Thread



Process and Thread

The task is to build a house



What do you need?




Process and Thread

The task is to build a house



What do you need?

Project




Process and Thread

The task is to build a house



What do you need?

Project



Materials
and instruments




Process and Thread

The task is to build a house



What do you need?

Project



Materials
and instruments



Land



Process and Thread

The task is to build a house




Workers



What do you need?

Project



Materials
and instruments



Land



Process and Thread


The task is to build a house



What do you need?

Project = Program

All of these is Process



Workers = Threads



Materials
and instruments = Data



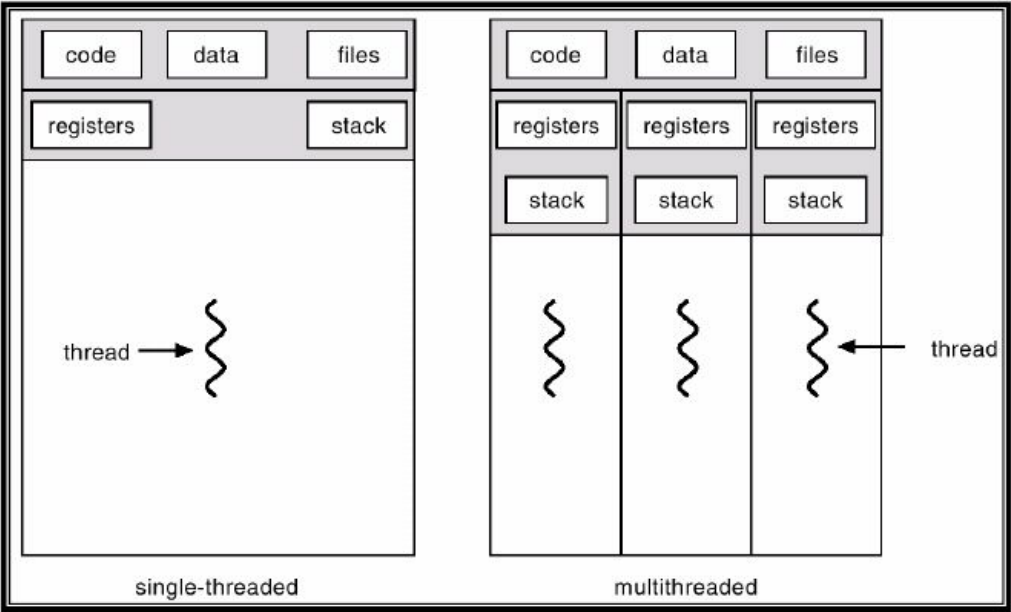
Land = Address space



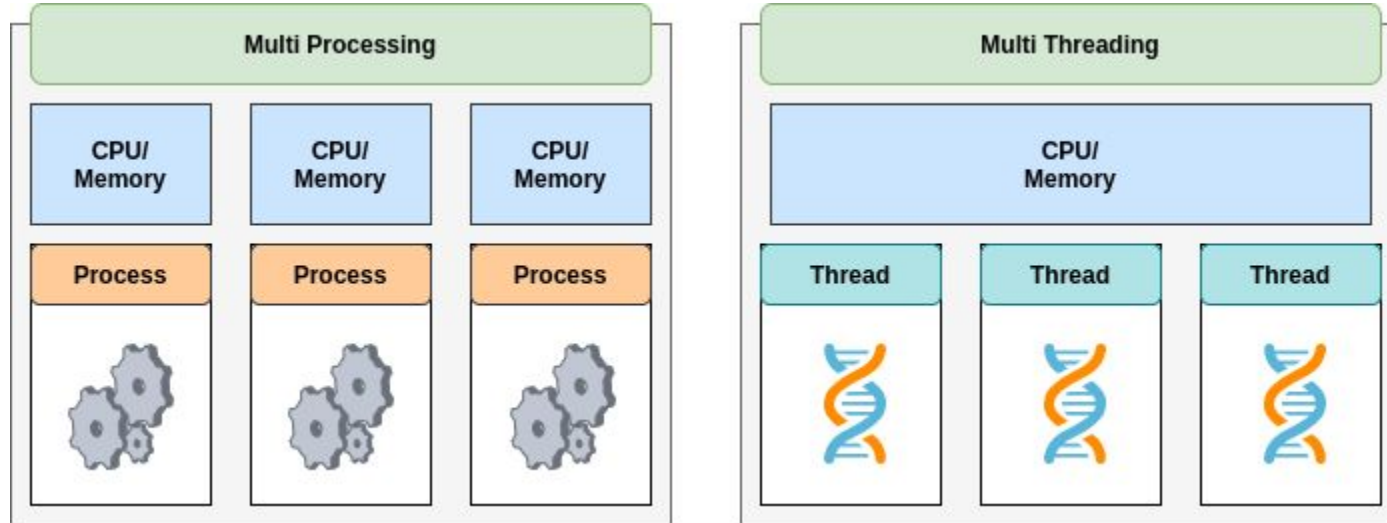
Process and Thread

Process – is a “container” for threads

Processes are isolated from each other, and threads can run in the same process.



Multi Threading



Multi-threaded programming is only possible on shared memory systems.

Algorithm?

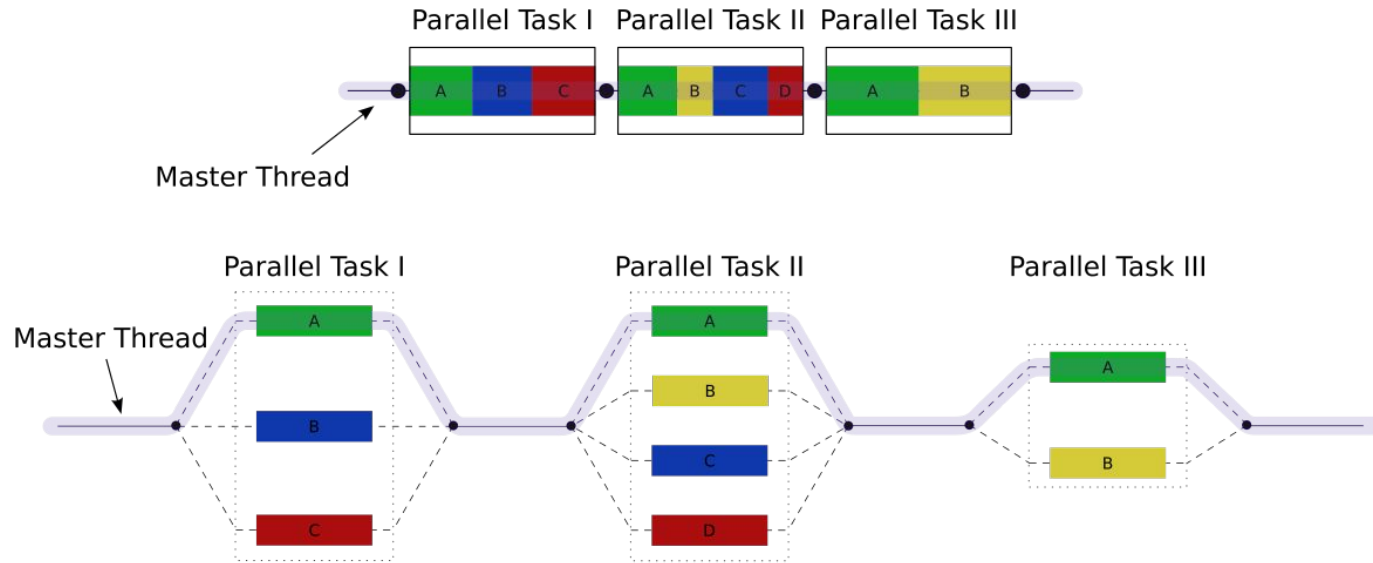
- many algorithms can be implemented with parallel computing
- anything we work with - is a data
- it's more about approaches to parallel computing than specific algorithms

Model of parallel computations proposed by Google.

- main idea is to process huge amount of data with thousands machines
- task is splitted on stages:
 - Map - apply simple instruction for each data sample on a local storages
 - Shuffle - reorder data based on keys constructed during map operation in order to keep all data with the same key on a same node
 - Reduce - process each group of reordered data

OpenMP

Fork – join parallelism



One version of the program for parallel and sequential execution.


SPMD (Single Program Multiple Data) model of parallel programming: the same code is used for all parallel threads.

OpenMP includes:

- compiler directives
- helper function
- environment variables

Basics directives embed directly into serial code

#pragma omp *directive-name* (clause(,) clause) ...)



```
...  
#pragma omp parallel  
{  
  
#pragma omp for  
for (int i = 0; i < 9 ; ++i) {  
    a[i] = i;  
}  
  
}  
...
```

```
#include <omp.h>
#include <iostream>
using namespace std;
int main()
{
    int a[100],b[100],c[100];

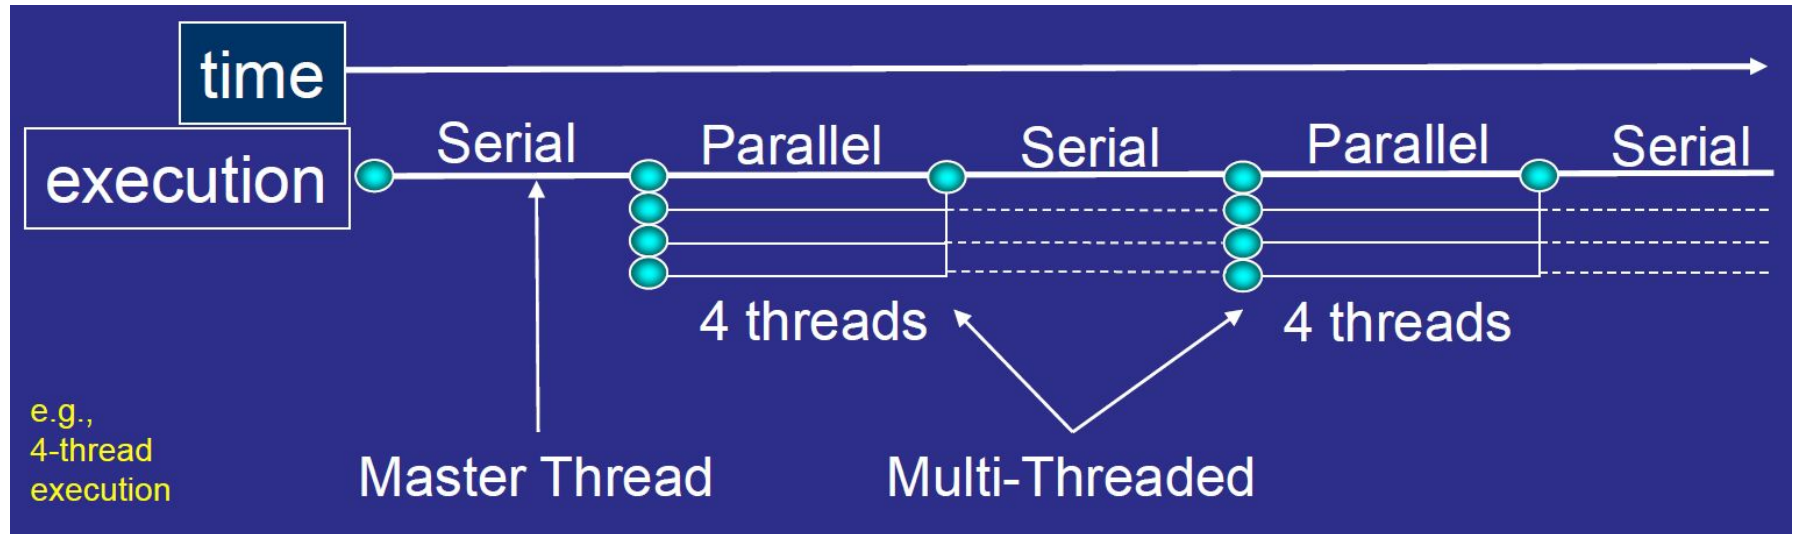
    for (int i=0;i<100;i++)
    {
        a[i]=1;
        b[i]=1;
    }

    for (int i=0;i<100;i++)
    {
        c[i]=a[i]+b[i];
    }
}
```

```
#include <omp.h>
#include <iostream>
using namespace std;
int main()
{
    int a[100],b[100],c[100];
    #pragma omp parallel for
    for (int i=0;i<100;i++)
    {
        a[i]=1;
        b[i]=1;
    }
    #pragma omp parallel for
    for (int i=0;i<100;i++)
    {
        c[i]=a[i]+b[i];
    }
}
```

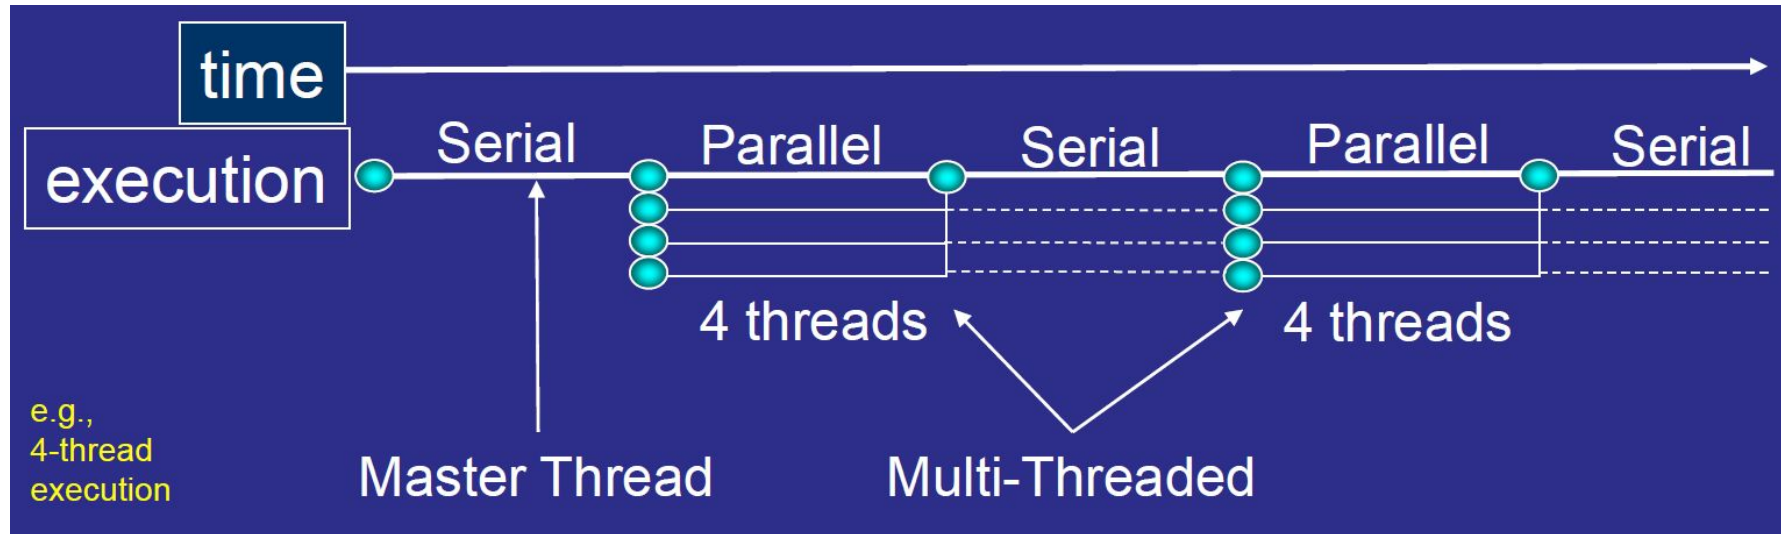

When the program starts, the main thread or master is spawned.

Only the main thread executes all sequential blocks of the program.



For create additional threads for parallel section you should write directive

```
#pragma omp parallel (clause(), clause...)
{
    \\ parallel code
}
```



For create additional threads for parallel section you should write directive

```
#pragma omp parallel (clause(), clause...)
{
    \\ parallel code
}
```

```
1  #include <stdio.h>
2  #include <locale>
3
4  int main()
5  {
6      printf("Serial block 1\n");
7      #pragma omp parallel
8      {
9          printf("Parallel block\n");
10     }
11     printf("Serial block 2\n");
12 }
```

For create additional threads for parallel section you should write directive

```
#pragma omp parallel num_threads(5)
```

```
#pragma omp parallel if (condition)
```

```
1  #include <iostream>
2  #include "omp.h"
3  #include <string>
4
5  using namespace std;
6  int main()
7  {
8      string hw = "Hello, world\n";
9      #pragma omp parallel num_threads(5)
10     {
11         cout << hw;
12     }
13     return 0;
14 }
```

Shared

Data race

```
1  #include <iostream>
2  #include "omp.h"
3  #include <string>
4
5  using namespace std;
6
7  int main()
8  {
9      int x = 0;
10     #pragma omp parallel shared(x) num_threads(30)
11     {
12         x += 1;
13     }
14     cout << "x = " << x << endl;
15     return 0;
16 }
```

Shared

Private

Data race

```
1  #include <iostream>
2  #include "omp.h"
3  #include <string>
4
5  using namespace std;
6
7  int main()
8  {
9      int x = 0;
10     #pragma omp parallel shared(x) num_threads(30)
11     {
12         x += 1;
13     }
14     cout << "x = " << x << endl;
15     return 0;
16 }
```

```
1  #include <stdio.h>
2  #include <omp.h>
3  #include <locale>
4
5  int main()
6  {
7      int n = 1;
8
9      printf("n in sequential area (start): %d\n", n);
10
11     #pragma omp parallel private(n) num_threads(4)
12     {
13         printf("The value of n in the thread (at the input): %d\n", n);
14
15         n = omp_get_thread_num(); // We assign n the number of the current thread
16         printf("The value of n in the thread (at the output): %d\n", n);
17     }
18     printf("n in sequential area (end): %d\n", n);
19     return 0;
20 }
```

private – create local variable for each thread

firstprivate – create local variable for each thread with initialization from previous serial part

```
1  #include <stdio.h>
2  #include <omp.h>
3  #include <locale>
4
5  int main()
6  {
7      int n = 1;
8
9      printf("The value of n at the beginning: %d\n", n);
10 #pragma omp parallel firstprivate(n)
11 {
12     printf("The value of n in the thread (at the input): %d\n", n);
13     n = omp_get_thread_num(); // assign the variable n to the sequence number of the thread
14     printf("The value of n in the thread (at the output): %d\n", n);
15 }
16 printf("The value of n at the end: %d\n", n);
17 }
```

private – create local variable for each thread

lastprivate – create variable after parallel part with initialization from last parallel section

```
1  #include <stdio.h>
2  #include <omp.h>
3  #include <locale>
4
5  int main()
6  {
7      int n = 1;
8      int i = 0;
9      int a;
10     printf("The value of n at the beginning: %d\n", n);
11     #pragma omp parallel for private(i) lastprivate(a) num_threads(5)
12     for(i=0;i<5;i++)
13     {
14         a=i+1;
15         n = omp_get_thread_num(); // assign the variable n to the sequence number of the thread
16         printf("The value of a in the thread %d\n: %d\n",a,n);
17     }
18     printf("The value of a at the end: %d\n", a);
19 }
```


single

- private
- firstprivate
- copyprivate
- nowait

```
1  #include <iostream>
2  #include <stdio.h>
3  #include <omp.h>
4  int main()
5  {
6      int num;
7      #pragma omp parallel num_threads(4) private(num)
8      {
9          num = omp_get_thread_num();
10         printf("Before the directive single num=%d \n", num);
11     #pragma omp barrier
12     #pragma omp single copyprivate(num)
13     {
14         printf("Enter an integer: ");
15         scanf("%d", &num);
16     }
17     printf("After the directive single num=%d \n", num);
18 }
19 }
```

single nowait

```
1  #include <iostream>
2  #include <stdio.h>
3  #include <omp.h>
4  int main()
5  {
6      double k = 0;
7      #pragma omp parallel num_threads(4) firstprivate(k)
8      {
9          printf("Before single without nowait \n");
10         #pragma omp single
11         {
12             // This loop is added so that the thread can do some work.
13             for (int i = 0; i < 100000; i++)
14             {
15                 k += (double)i / (i + 1);
16             }
17             printf("In single directive\n");
18         }
19         printf("After the single directive without nowait. This message will never be earlier than the previous ones. k = %f \n",k);
20
21         #pragma omp barrier // This directive synchronizes threads
22
23         printf("Before single directive with nowait \n");
24         #pragma omp single nowait
25         {
26             // his loop is added so that the thread can do some work.
27             for (int i = 0; i < 100000; i++)
28             {
29                 k += (double)i / (i + 1);
30             }
31             printf("In single directive\n");
32         }
33         printf("After the single directive with nowait. This message may be earlier than the previous ones. k = %f \n", k);
34     }
35     return 0;
36 }
```

master

```
1  #include <stdio.h>
2  #include <locale>
3  #include <omp.h>
4
5  int main()
6  {
7
8      int n;
9      #pragma omp parallel private(n)
10     {
11         n = 1;
12         #pragma omp master
13         {
14             n = 2;
15         }
16         printf("The first value of the n thread %d: %d\n", omp_get_thread_num(), n);
17         #pragma omp barrier
18         #pragma omp master
19         {
20             n = 3;
21         }
22         printf("The second value of the n thread %d: %d\n", omp_get_thread_num(), n);
23     }
24     return 0;
25 }
```

Parallelizing loops



```
1  #include <iostream>
2  #include <omp.h>
3  using namespace std;
4  int main()
5  {
6      const long int n = 40000000;
7      double* a = new double[n];
8      double* b = new double[n];
9      double* c = new double[n];
10     for (long int i = 0; i < n; i++)
11     {
12         a[i] = (double)rand() / RAND_MAX;
13         b[i] = (double)rand() / RAND_MAX;
14     }
15     double time = omp_get_wtime();
16
17     for (long int i = 0; i < n; i++)
18     {
19         c[i] = a[i] + b[i];
20     }
21     cout << "c[100]=" << c[100] << endl;
22     cout << "Time = " << (omp_get_wtime() - time) << endl;
23     delete[] a;
24     delete[] b;
25     delete[] c;
26 }
```

Without omp

```
1  #include <iostream>
2  #include <omp.h>
3  using namespace std;
4  int main()
5  {
6      const Long int n = 40000000;
7      double* a = new double[n];
8      double* b = new double[n];
9      double* c = new double[n];
10     for (Long int i = 0; i < n; i++)
11     {
12         a[i] = (double)rand() / RAND_MAX;
13         b[i] = (double)rand() / RAND_MAX;
14     }
15     double time = omp_get_wtime();
16
17     for (Long int i = 0; i < n; i++)
18     {
19         c[i] = a[i] + b[i];
20     }
21     cout << "c[100]=" << c[100] << endl;
22     cout << "Time = " << (omp_get_wtime() - time) << endl;
23     delete[] a;
24     delete[] b;
25     delete[] c;
26 }
```

With omp

```
1  #include <iostream>
2  #include <omp.h>
3  using namespace std;
4  int main()
5  {
6      const Long int n = 40000000;
7      double* a = new double[n];
8      double* b = new double[n];
9      double* c = new double[n];
10     for (Long int i = 0; i < n; i++)
11     {
12         a[i] = (double)rand() / RAND_MAX;
13         b[i] = (double)rand() / RAND_MAX;
14     }
15     double time = omp_get_wtime();
16     #pragma omp parallel shared(a,b,c)
17     for (Long int i = 0; i < n; i++)
18     {
19         c[i] = a[i] + b[i];
20     }
21     cout << "c[100]=" << c[100] << endl;
22     cout << "Time = " << (omp_get_wtime() - time) << endl;
23     delete[] a;
24     delete[] b;
25     delete[] c;
26 }
```

With omp, parallel for

```
1  #include <iostream>
2  #include <omp.h>
3  using namespace std;
4  int main()
5  {
6      const long int n = 40000000;
7      double* a = new double[n];
8      double* b = new double[n];
9      double* c = new double[n];
10     for (long int i = 0; i < n; i++)
11     {
12         a[i] = (double)rand() / RAND_MAX;
13         b[i] = (double)rand() / RAND_MAX;
14     }
15     double time = omp_get_wtime();
16     #pragma omp parallel shared(a,b,c)
17     for (long int i = 0; i < n; i++)
18     {
19         c[i] = a[i] + b[i];
20     }
21     cout << "c[100]=" << c[100] << endl;
22     cout << "Time = " << (omp_get_wtime() - time) << endl;
23     delete[] a;
24     delete[] b;
25     delete[] c;
26 }
```

```
1  #include <iostream>
2  #include <omp.h>
3  using namespace std;
4  int main()
5  {
6      const long int n = 40000000;
7      double* a = new double[n];
8      double* b = new double[n];
9      double* c = new double[n];
10     for (long int i = 0; i < n; i++)
11     {
12         a[i] = (double)rand() / RAND_MAX;
13         b[i] = (double)rand() / RAND_MAX;
14     }
15     double time = omp_get_wtime();
16     for (int j = 0; j < 100; j++)
17     {
18         #pragma omp parallel shared(a,b,c)
19         #pragma omp for
20         for (long int i = 0; i < n; i++)
21         {
22             c[i] = a[i] + b[i];
23         }
24     }
25     cout << "c[100]=" << c[100] << endl;
26     cout << "Time = " << (omp_get_wtime() - time) / 100 << endl;
27     delete[] a;
28     delete[] b;
29     delete[] c;
30 }
```